

# Preparing the environment

Controlling the presence of GNU C Compiler

```
gbiondo@LinuxMint:~/Development$ gcc -v
Using built-in specs.
COLLECT_GCC=gcc
COLLECT_LTO_WRAPPER=/usr/lib/gcc/x86_64-linux-gnu/7/lto-wrapper
OFFLOAD_TARGET_NAMES=nvptx-none
OFFLOAD_TARGET_DEFAULT=1
Target: x86_64-linux-gnu
Configured with: ../src/configure -v --with-pkgversion='Ubuntu 7.4.0-1ubuntu1~18.04.1'
--with-bugurl=file:///usr/share/doc/gcc-7/README.Bugs --enable-languages=c,ada,c++,go,brig,d,fortran,objc,obj-c++ --prefix=/usr --with-gcc-major-version-only --program-suffix=-7 --program-prefix=x86_64-linux-gnu- --enable-shared --enable-linker-build-id --libexecdir=/usr/lib --without-included-gettext --enable-threads=posix --libdir=/usr/lib --enable-nls --with-sysroot=/ --enable-clocale=gnu --enable-libstdcxx-debug --enable-libstdcxx-time=yes --with-default-libstdcxx-abi=new --enable-gnu-unique-object --disable-vtable-verify --enable-libmpx --enable-plugin --enable-default-pie --with-system-zlib --with-target-system-zlib --enable-objc-gc=auto --enable-multiarch --disable-werror --with-arch-32=i686 --with-abi=m64 --with-multilib-list=m32,m64,mx32 --enable-multilib --with-tune=generic --enable-offload-targets=nvptx-none --without-cuda-driver --enable-checking=release --build=x86_64-linux-gnu --host=x86_64-linux-gnu --target=x86_64-linux-gnu
Thread model: posix
gcc version 7.4.0 (Ubuntu 7.4.0-1ubuntu1~18.04.1)
gbiondo@LinuxMint:~/Development$
```

Controlling the presence of Python

```
gbiondo@LinuxMint:~/Development$ python -V
Python 2.7.17
gbiondo@LinuxMint:~/Development$
```

Controlling the presence of nasm

```
gbiondo@LinuxMint:~/Development$ nasm -v
NASM version 2.13.02
gbiondo@LinuxMint:~/Development$
```

Controlling the presence of ld

```
gbiondo@LinuxMint:~/Development$ ld -v
GNU ld (GNU Binutils for Ubuntu) 2.30
gbiondo@LinuxMint:~/Development$
```

# Registers

A CPU register, or just register, is a temporary storage or working location built into the CPU itself, separate from memory. Computations are typically performed by the CPU using registers. It is a quickly accessible location available to a computer's central processing unit (CPU). Registers usually consist of a small amount of fast storage, although some registers have specific hardware functions, and may be read-only or write-only.

X64 assembly code uses sixteen 64-bit registers. Additionally, the lower bytes of some of these registers may be accessed independently as 32-, 16- or 8-bit registers. The register names are as follows:

| Bytes 0-7<br>(64 bits) | Bytes 0-3<br>(32 bits) | Bytes 0-1<br>(16 bits) | Byte 0<br>(8 bits) | Usage  |
|------------------------|------------------------|------------------------|--------------------|--|
| rax                    | eax                    | ax                     | al                 | Used in arithmetic operations  |
| rcx                    | ecx                    | cx                     | cl                 | Used in shift/rotate instructions and loops  |
| rdx                    | edx                    | dx                     | dl                 | Used in arithmetic operations and I/O operations   |
| rbx                    | ebx                    | bx                     | bl                 | Used as a pointer to data (located in segment register DS, when in segmented mode)   |
| rsi                    | esi                    | si                     | sil                | Used as a pointer to a source in stream operations   |
| rdi                    | edi                    | di                     | dil                | Used as a pointer to a destination in stream operations  |
| rsp                    | esp                    | sp                     | spl                | Pointer to the top of the stack, and should not be used for data or other uses.  |
| rbp                    | ebp                    | bp                     | bpl                | Used to point to the base of the stack. It is used as a base pointer during function calls. This register should not be used for data or other uses. |
| r8                     | r8d                    | r8w                    | r8b                |  |
| r9                     | r9d                    | r9w                    | r9b                |  |
| r10                    | r10d                   | r10w                   | r10b               |  |
| r11                    | r11d                   | r11w                   | r11b               |  |
| r12                    | r12d                   | r12w                   | r12b               |  |
| r13                    | r13d                   | r13w                   | r13b               |  |
| r14                    | r14d                   | r14w                   | r14b               |  |
| r15                    | r15d                   | r15w                   | r15b               |  |

---

## Instruction Pointer Register (RIP)

In addition to the GPRs, there is a special register, **RIP**, which is used by the CPU to point to the next instruction to be executed. Specifically, since the **RIP** points to the next instruction, that

means the instruction being pointed to by **RIP**, and shown in the debugger, has not yet been executed.

---

## Flag Register (rFlags)

The flag register, **rFlags**, is used for status and CPU control information. The **rFlags** register is updated by the CPU after each instruction and not directly accessible by programs. This register stores status information about the instruction that was just executed. Of the 64-bits in the rFlag register, many are reserved for future use. We will come back to this topic later in greater detail.

## Access modes

All registers can be accessed in 16-bit and 32-bit modes. In 16-bit mode, the register is identified by its two-letter abbreviation from the list above. In 32-bit mode, this two-letter abbreviation is prefixed with an 'E' (extended). For example, 'EAX' is the accumulator register as a 32-bit value. Similarly, in the 64-bit version, the 'E' is replaced with an 'R' (register), so the 64-bit version of 'EAX' is called 'RAX'.

It is also possible to address the first four registers (AX, CX, DX and BX) in their size of 16-bit as two 8-bit halves. The least significant byte (LSB), or low half, is identified by replacing the 'X' with an 'L'. The most significant byte (MSB), or high half, uses an 'H' instead. For example, CL is the LSB of the counter register, whereas CH is its MSB.

In total, this gives us five ways to access the accumulator, counter, data and base registers: 64-bit, 32-bit, 16-bit, 8-bit LSB, and 8-bit MSB. The other four are accessed in only four ways: 64-bit, 32-bit, 16-bit, and 8-bit.

|         | ACCUMULATOR |  | COUNTER |  | DATA  |  | BASE  |  | STACK PTR |  | STACK BASE |  | SOURCE  |  | DESTINATION |  |
|---------|-------------|--|---------|--|-------|--|-------|--|-----------|--|------------|--|---------|--|-------------|--|
| 64 BITS | RAX         |  | RCX     |  | RDX   |  | RBX   |  | RSP       |  | RBP        |  | RSI     |  | RDI         |  |
| 32 BITS | EAX         |  | ECX     |  | EDX   |  | EBX   |  | ESP       |  | EBP        |  | ESI     |  | EDI         |  |
| 16 BITS | AX          |  | CX      |  | DX    |  | BX    |  | SP        |  | BP         |  | SI      |  | DI          |  |
| 8 BITS  | AH AL       |  | CH CL   |  | DH DL |  | BH BL |  | SPH SPL   |  | BPH BPL    |  | SIH SIL |  | DIH DIL     |  |

---

## Data Types

The following table explains the data types in assembly based on length:

| Name       | Directive | Bytes |
|------------|-----------|-------|
| Byte       | db        | 1     |
| Word       | dw        | 2     |
| Doubleword | dd        | 4     |
| Quadword   | dq        | 8     |

# First steps in assembly programming

## The format of a program

We need first the system call (syscall) `exit`. To find it, we can look at the file

`/usr/include/x86_64-linux-gnu/asm/unistd_64.h`

as follows

```
gbiondo@LinuxMint:~/Development$ cat /usr/include/x86_64-linux-gnu/asm/unistd_64.h |  
grep -i exit  
#define __NR_exit 60  
#define __NR_exit_group 231  
gbiondo@LinuxMint:~/Development$
```

Once this is known (the `syscall` is the number `60`), we need to investigate how it works and the parameters it requires:

```
gbiondo@LinuxMint:~/Development$ man 2 exit  
NAME  
    _exit, _Exit - terminate the calling process  
  
SYNOPSIS  
    #include <unistd.h>  
  
    void _exit(int status);  
  
    #include <stdlib.h>  
  
    void _Exit(int status);  
  
    Feature Test Macro Requirements for glibc (see feature_test_macros(7)):  
    ...  
gbiondo@LinuxMint:~/Development$
```

Chiefly, this instruction only requires one parameter, `status`, and returns nothing (it is a `void` function).

Traditionally, the first program one writes is the well known 'hello world'. Let's stick to this tradition - so we only need to know how to write on the screen. We know how to do this; first we look for the `write` syscall:

```
gbiondo@LinuxMint:~/Development$ cat /usr/include/x86_64-linux-gnu/asm/unistd_64.h |  
grep -i write  
#define __NR_write 1  
#define __NR_pwrite64 18  
#define __NR_writev 20  
#define __NR_pwritev 296  
#define __NR_process_vm_writev 311  
#define __NR_pwrite2 328  
gbiondo@LinuxMint:~/Development$
```

which has the number `1` associated; then we look at its man page:

```

gbiondo@LinuxMint:~/Development$ man 2 write
NAME
    write - write to a file descriptor

SYNOPSIS
    #include <unistd.h>

    ssize_t write(int fd, const void *buf, size_t count);

DESCRIPTION
    write() writes up to count bytes from the buffer starting at buf to the file
    referred to by the file descriptor fd.
...
gbiondo@LinuxMint:~/Development$

```

The system call has then three arguments: a file descriptor, the location of memory containing the value to print, and the number of bytes to print.

File descriptors are usual Unix ones:

| Value    | Name            | Alias in <code>stdio.h</code> |
|----------|-----------------|-------------------------------|
| <b>0</b> | Standard Input  | stdin                         |
| <b>1</b> | Standard Output | stdout                        |
| <b>2</b> | Standard Error  | stderr                        |

In order to invoke a syscall, the programmer should load the RAX register with the syscall number (in this case, 1), and the arguments must be load as follows:

| Argument | Register |
|----------|----------|
| <b>1</b> | RDI      |
| <b>2</b> | RSI      |
| <b>3</b> | RDX      |
| <b>4</b> | R10      |
| <b>5</b> | R8       |
| <b>6</b> | R9       |

In this case, we will proceed by loading the registers as follows:

| Register   | VALUE                    |
|------------|--------------------------|
| <b>RAX</b> | 1 (syscall number)       |
| <b>RDI</b> | 1 (stdout)               |
| <b>RSI</b> | Address of "Hello world" |

| Register | VALUE                   |
|----------|-------------------------|
| RDX      | Length of "Hello world" |

A variable will be defined, `hello_world`, containing the string `'hello world'`, terminated by a line feed (`\n`), whose ASCII encoding is `10`, or `0xa`. This is a string of bytes, so the data type will be `db`.

Another variable to be defined is the length of the `hello_world` string. This is obtained by the instruction `$-`, which evaluates the current line.

The two variables will be declared in the `.data` section of our program. The `.text` section, containing the real program, will just need to load the registers as previously described and to invoke the system call.

The code is as follows:

```
global _start
section .text
_start:
    mov rax, 1
    mov rdi, 1
    mov rsi, hello_world
    mov rdx, length
    syscall

section .data
    hello_world: db 'hello world',0xa
    length: equ $-hello_world
```

We compile, link and run the executable as follows:

```
gbiondo@LinuxMint:~/Development/HelloWorld$ nasm -felf64 hello.nasm -o hello.o
gbiondo@LinuxMint:~/Development/HelloWorld$ ld hello.o -o helloWorld
gbiondo@LinuxMint:~/Development/HelloWorld$ ./helloWorld
hello world
Segmentation fault (core dumped)
gbiondo@LinuxMint:~/Development/HelloWorld$ echo $?
139
gbiondo@LinuxMint:~/Development/HelloWorld$
```

The core dump: Unix systems return error n.  $128 + \text{signal}$  when a signal is received.  $128 + 11 = 139$ . Signal 11 is SIGSEV (i.e. segmentation violation). In fact, there has been a segmentation violation, in the sense that the program invokes no `'exit'` syscall. In fact, changing the program as follows:

```

global _start
section .text
_start:
    mov rax, 1
    mov rdi, 1
    mov rsi, hello_world
    mov rdx, length
    syscall

    mov rax, 0x3C
    syscall

section .data
    hello_world: db 'hello world',0xa
    length: equ $-hello_world

```

solves the problem.

## Fun with Arithmetics

The very first program does not do so much:

```

global _start
section .text
_start:
    mov rax,0x1
    add rax,0x2

    mov rax, 60
    mov rdi, 0
    syscall

```

An analysis with GDB can help further. The program has been compiled as follows:

```

gbiondo@LinuxMint:~/Development/sumDiffs $ nasm -felf64 main1.nasm -g -F dwarf -o
main1a1.o
gbiondo@LinuxMint:~/Development/sumDiffs $ ld main1a1.o -o main1a1

```

And we invoke GDB as follows:

```

gbiondo@LinuxMint:~/Development/sumDiffs $ gdb main1a1
GNU gdb (Ubuntu 8.1-0ubuntu3.2) 8.1.0.20180409-git
Copyright (C) 2018 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.
For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from main1a1...done.

```

First, we want to see the source code, to understand where putting breakpoints

```

(gdb) disassemble _start
Dump of assembler code for function _start:
   0x0000000000400080 <+0>:      mov     $0x1,%eax
   0x0000000000400085 <+5>:      add     $0x2,%rax
   0x0000000000400089 <+9>:      mov     $0x3c,%eax
   0x000000000040008e <+14>:     mov     $0x0,%edi
   0x0000000000400093 <+19>:     syscall
End of assembler dump.
(gdb)

```

It comes handy to put a breakpoint on the first line, to see the initial status; after the first mov instruction, and after the add instruction:

```

(gdb) break _start
Breakpoint 1 at 0x400080: file main1.nasm, line 7.
(gdb) break 8
Breakpoint 2 at 0x400085: file main1.nasm, line 8.
(gdb) break 9
Breakpoint 3 at 0x400089: file main1.nasm, line 9.
(gdb) info breakpoints
Num      Type             Disp Enb Address                  What
1        breakpoint      keep y  0x0000000000400080 main1.nasm:7
2        breakpoint      keep y  0x0000000000400085 main1.nasm:8
3        breakpoint      keep y  0x0000000000400089 main1.nasm:9

```

Starting the program and checking the **EAX** register, which should contain the default value (which is zero):

```

(gdb) run
Starting program: /home/gbiondo/Development/sumsDiffs/main1a1

Breakpoint 1, _start () at main1.nasm:7
7      mov rax,0x1
(gdb) info registers eax
eax                0x0  0
(gdb) info registers rax
rax                0x0  0

```

Taking a step, the first instruction gets executed, so the register is loaded with the value 1:



```
(gdb) s
Breakpoint 2, _start () at main1.nasm:8
8      add rax,0x2
(gdb) info registers eax
eax      0x1  1
(gdb) info registers rax
rax      0x1  1
```

We then run the second instruction, adding 2 to the contents of the **RAX** registry, we obtain:

```
(gdb) info registers eax
eax      0x3  3
(gdb) info registers rax
rax      0x3  3
(gdb)
```

This shows also how the **EAX** and **RAX** are interchangeable, in this context.

Let us add something more to the code to see some further functionalities. Consider the following:

```
global _start
section .text
_start:
    mov rax,0x1
    add rax,0x2

    mov rbx, rax
    add rbx, rbx
    add bl, byte [addr1]

    mov rax, 60
    mov rdi, 0
    syscall

section .data
    addr1: db 0x23
```

(keep in mind that 0x23 = 35 in decimal)

The analysis with GDB is quite interesting. As usual we disassemble the code and add some breakpoints:

```

gbiondo@LinuxMint:~/Development/sumDiffs $ gdb main2
GNU gdb (Ubuntu 8.1-0ubuntu3.2) 8.1.0.20180409-git
Copyright (C) 2018 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.
For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from main2...done.
(gdb) set disassembly-flavor intel
(gdb) disassemble _start
Dump of assembler code for function _start:
   0x00000000004000b0 <+0>:      mov     eax,0x1
   0x00000000004000b5 <+5>:      add     rax,0x2
   0x00000000004000b9 <+9>:      mov     rbx,rax
   0x00000000004000bc <+12>:     add     rbx,rbx
   0x00000000004000bf <+15>:     add     bl,BYTE PTR ds:0x6000d4
   0x00000000004000c6 <+22>:     mov     eax,0x3c
   0x00000000004000cb <+27>:     mov     edi,0x0
   0x00000000004000d0 <+32>:     syscall
End of assembler dump.

(gdb) info breakpoints
Num      Type             Disp Enb Address                  What
1        breakpoint       keep y  0x00000000004000b9  main2.nasm:9
2        breakpoint       keep y  0x00000000004000bc  main2.nasm:11
3        breakpoint       keep y  0x00000000004000c6  main2.nasm:13
4        breakpoint       keep y  0x00000000004000bf  main2.nasm:12
(gdb) run

```

We first move the value of **RAX** into **RBX**, hence **RBX** should have a value of 3; then we add to **RBX** the value of **RBX**, which should give us a result of 6; finally we add to **RBX** the contents of the memory address **0x6000d4**, which is the address of **addr1**. The result should be 41 (6 + 35).

We have:

```

Breakpoint 1, _start () at main2.nasm:10
10      mov     rbx, rax
(gdb) i r rbx
rbx      0x0  0
(gdb) s

Breakpoint 2, _start () at main2.nasm:11
11      add     rbx, rbx
(gdb) i r rbx
rbx      0x3  3
(gdb) s

Breakpoint 4, _start () at main2.nasm:12
12      add     bl, byte [addr1]
(gdb) i r rbx
rbx      0x6  6
(gdb) s

Breakpoint 3, _start () at main2.nasm:14
14      mov     rax, 60
(gdb) i r rbx
rbx      0x29 41

```

And results, here, are in line with the expectations.

Consider the following code, instead:

```
global _start
section .text
_start:
    mov cl, byte [addr2]
    sub cl, byte [addr1]

    mov rax, 60
    mov rdi, 0
    syscall

section .data
    addr1: db 0x23
    addr2: db 0x17
```

GDB'ing it:

```
(gdb) set disassembly-flavor intel
(gdb) disassemble _start
Dump of assembler code for function _start:
   0x00000000004000b0 <+0>:    mov     cl,BYTE PTR ds:0x6000cd
   0x00000000004000b7 <+7>:    sub     cl,BYTE PTR ds:0x6000cc
   0x00000000004000be <+14>:   mov     eax,0x3c
   0x00000000004000c3 <+19>:   mov     edi,0x0
   0x00000000004000c8 <+24>:   syscall
End of assembler dump.
(gdb) break 8
Breakpoint 1 at 0x4000b7: file main3.nasm, line 8.
(gdb) break 9
Breakpoint 2 at 0x4000be: file main3.nasm, line 9.
(gdb) info breakpoints
Num      Type             Disp Enb Address            What
1        breakpoint       keep y   0x00000000004000b7 main3.nasm:8
2        breakpoint       keep y   0x00000000004000be main3.nasm:9
(gdb) run
Starting program: /home/gbiondo/Development/sumsDiffs/main3

Breakpoint 1, _start () at main3.nasm:8
8          sub cl, byte [addr1]
(gdb) i r cl
cl          0x17 23
(gdb) s

Breakpoint 2, _start () at main3.nasm:10
10         mov rax, 60
(gdb) i r cl
cl          0xf4 -12
(gdb)
```

Results are obviously in line with expectations. Interestingly, we worked only with the 'small registers'.

## Flags

The FLAGS register is the status register in Intel x86 microprocessors that contains the current state of the processor. This register is 16 bits wide. Its successors, the EFLAGS and RFLAGS registers, are 32 bits and 64 bits wide, respectively. The wider registers retain compatibility with their smaller predecessors.

The fixed bits at bit positions 1, 3 and 5, and carry, parity, adjust, zero and sign flags.

The flags register is as follows:

| Bit #         | Mask   | Abbreviation | Description           | Notes   |
|---------------|--------|--------------|-----------------------|---|
| <b>0</b>      | 0x0001 | CF           | Carry flag            | Set if the last arithmetic operation carried (addition) or borrowed (subtraction) a bit beyond the size of the register. This is then checked when the operation is followed with an add-with-carry or subtract-with-borrow to deal with values too large for just one register to contain. |
| <b>1</b>      | 0x0002 |              | Reserved              |   |
| <b>2</b>      | 0x0004 | PF           | Parity flag           | Set if the number of set bits in the least significant byte is a multiple of 2.   |
| <b>3</b>      | 0x0008 |              | Reserved              |   |
| <b>4</b>      | 0x0010 | AF           | Adjust flag           | Carry of Binary Code Decimal (BCD) numbers arithmetic operations.   |
| <b>5</b>      | 0x0020 |              | Reserved              |   |
| <b>6</b>      | 0x0040 | ZF           | Zero flag             | Set if the result of an operation is Zero (0).  |
| <b>7</b>      | 0x0080 | SF           | Sign flag             | Set if the result of an operation is negative.  |
| <b>8</b>      | 0x0100 | TF           | Trap flag             | Set if step by step debugging.  |
| <b>9</b>      | 0x0200 | IF           | Interrupt enable flag | Set if interrupts are enabled.  |
| <b>10</b>     | 0x0400 | DF           | Direction flag        | If set, string operations will decrement their pointer rather than incrementing it, reading memory backwards.   |
| <b>11</b>     | 0x0800 | OF           | Overflow flag         | Set if signed arithmetic operations result in a value too large for the register to contain.  |
| <b>12, 13</b> | 0x3000 | IOPL         | I/O privilege lvl.    | I/O Privilege Level of the current process.   |
| <b>14</b>     | 0x4000 | NT           | Nested task flag      | Controls chaining of interrupts. Set if the current process is linked to the next process.  |
| <b>15</b>     | 0x8000 |              | Reserved              |   |

In 32 and 64-bits architectures, the following flags are available as well:

| Bit #     | Abbreviation | Description | Notes                         |
|-----------|--------------|-------------|-------------------------------|
| <b>16</b> | RF           | Resume Flag | Response to debug exceptions. |

| Bit # | Abbreviation | Description                     | Notes   |
|-------|--------------|---------------------------------|---|
| 17    | VM           | Virtual-8086 Mode               | Set if in 8086 compatibility mode.                      |
| 18    | AC           | Alignment Check                 | Set if alignment checking of memory references is done. |
| 19    | VIF          | Virtual Interrupt Flag          | Virtual image of IF.                                    |
| 20    | VIP          | Virtual Interrupt Pending flag. | Set if an interrupt is pending.                         |
| 21    | ID           | Identification Flag             | Support for CPUID instruction if can be set.            |

Consider the following piece of code:

```
global _start
section .text
_start:
    mov rax, 0x0
    stc
    adc rax, 0x0
    stc
    adc rax, 0x0
    stc
    stc

    mov rax, 60
    mov rdi, 0
    syscall

section .data
```

We prepare for the execution as follows:

```
(gdb) disassemble _start
Dump of assembler code for function _start:
0x000000000400080 <+0>:      mov     eax,0x0
0x000000000400085 <+5>:      stc
0x000000000400086 <+6>:      adc     rax,0x0
0x00000000040008a <+10>:     stc
0x00000000040008b <+11>:     adc     rax,0x0
0x00000000040008f <+15>:     stc
0x000000000400090 <+16>:     stc
0x000000000400091 <+17>:     mov     eax,0x3c
0x000000000400096 <+22>:     mov     edi,0x0
0x00000000040009b <+27>:     syscall
(gdb) info breakpoints
Num    Type            Disp Enb Address            What
1      breakpoint      keep y  0x000000000400080 jumpingJack.nasm:7
2      breakpoint      keep y  0x000000000400085 jumpingJack.nasm:8
3      breakpoint      keep y  0x000000000400086 jumpingJack.nasm:9
4      breakpoint      keep y  0x00000000040008a jumpingJack.nasm:10
5      breakpoint      keep y  0x00000000040008b jumpingJack.nasm:11
6      breakpoint      keep y  0x00000000040008f jumpingJack.nasm:12
7      breakpoint      keep y  0x000000000400090 jumpingJack.nasm:13
8      breakpoint      keep y  0x000000000400091 jumpingJack.nasm:14
```

The following values have been obtained by invoking at every single breakpoints the following commands:

```
(gdb) i r eax
(gdb) p $eflags
(gdb) p/t $eflags
```

Finally, the instruction **STC** sets the Carry Flag, taking no operand; whereas in its most basic form, the instruction **ADC** takes a Register **R** and an Operand **O**, and performs the Add With Carry operation as follows:  $R = R + O + CF$ .

we have:

| Breakpoint | i r eax   | p \$eflags | p/t \$eflags |
|------------|-----------|------------|--------------|
| 1          | eax 0x0 0 | [ IF ]     | 1000000010   |
| 2          | eax 0x0 0 | [ IF ]     | 1000000010   |
| 3          | eax 0x0 0 | [ CF IF ]  | 1000000011   |
| 4          | eax 0x1 1 | [ IF ]     | 1000000010   |
| 5          | eax 0x1 1 | [ CF IF ]  | 1000000011   |
| 6          | eax 0x2 2 | [ IF ]     | 1000000010   |
| 7          | eax 0x2 2 | [ CF IF ]  | 1000000011   |
| 8          | eax 0x2 2 | [ CF IF ]  | 1000000011   |

Observe the following:

1. After an **ADC** operation, the **CF** is reset to 0
2. Two subsequent **STC** operation don't reset **CF** values.

## More fun with Arithmetics

Consider the following code:

```
(gdb) disassemble _start
Dump of assembler code for function _start:
0x000000000400080 <+0>:    mov     eax,0x17
0x000000000400085 <+5>:    stc
0x000000000400086 <+6>:    sbb     rax,0x22
0x00000000040008a <+10>:   mov     eax,0x3c
0x00000000040008f <+15>:   mov     edi,0x0
0x000000000400094 <+20>:   syscall
```

The instruction **SBB** is the Integer Subtraction with Borrow. In its most basic syntax, it takes a register (the destination, **DEST**) and subtracts from it the Operand (**SRC**) and the Carry Flag (**CF**) summed together:  $DEST = DEST - (SRC + CF)$ . In this case the expected result is 0x17 (which is 23) minus the sum of 0x22 (which is 34) and the carry flag (1). Shortly, we expect **RAX** to contain  $23 - (34 + 1) = -12$ .

Let's debug. The initial values are

```
(gdb) i r eax
eax          0x0  0
(gdb) p $eflags
$1 = [ IF ]
(gdb) p/t $eflags
$2 = 1000000010
```

After performing the initialisation of RAX, and after having set the carry flag, we have

```
(gdb) i r eax
eax          0x17 23
(gdb) p $eflags
$6 = [ CF IF ]
(gdb) p/t $eflags
$7 = 1000000011
```

Performing the subtraction gives the following results:

```
(gdb) i r eax
eax          0xffffffff4 -12
(gdb) p $eflags
$8 = [ CF SF IF ]
(gdb) p/t $eflags
$9 = 1010000011
```

which is interesting, as it shows:

- a) the way negative numbers are represented. We'll come back on this later.
- b) the carry flag is not affected by SBB
- c) the operation sets the Sign Flag (SF).

---

## Memory representation of negative numbers

In the last example, we saw first EAX loaded with `0x00000017`, and then with `0xFFFFFFFF4` to represent -12. What happened?

| Value      | Byte 7 | Byte 6 | Byte 5 | Byte 4 | Byte 3 | Byte 2 | Byte 1 | Byte 0 |
|------------|--------|--------|--------|--------|--------|--------|--------|--------|
| <b>17</b>  | 0      | 0      | 0      | 0      | 0      | 0      | 1      | 7      |
| <b>12</b>  | 0      | 0      | 0      | 0      | 0      | 0      | 1      | 2      |
| <b>-12</b> | F      | F      | F      | F      | F      | F      | F      | 4      |

The negative numbers are represented with the so called '2 complement'.

In binary, 12 is 00001100. First we calculate the 1-complement, which is obtained by swapping all 0s with 1s and vice versa. The result is quite straightforward.

---

## Increment and decrement values

Consider the following piece of code:

```
0x0000000000400080 <+0>: mov    eax,0x3b
0x0000000000400085 <+5>:      inc    rax
0x0000000000400088 <+8>:      inc    rax
0x000000000040008b <+11>:     inc    rax
0x000000000040008e <+14>:     dec    rax
0x0000000000400091 <+17>:     dec    rax
0x0000000000400094 <+20>:     mov    edi,0x0
0x0000000000400099 <+25>:     syscall
```

When debugging it:

```
(gdb) i breakpoints
Num    Type           Disp Enb Address            What
1      breakpoint     keep y   0x0000000000400085 incDec.asm:8
2      breakpoint     keep y   0x0000000000400088 incDec.asm:9
3      breakpoint     keep y   0x000000000040008b incDec.asm:10
4      breakpoint     keep y   0x000000000040008e incDec.asm:11
5      breakpoint     keep y   0x0000000000400091 incDec.asm:12
6      breakpoint     keep y   0x0000000000400094 incDec.asm:13
...
(gdb) r
Starting program: /home/gbiondo/Development/incsDecs/incDec

Breakpoint 1, _start () at incDec.asm:8
8          inc rax
(gdb) i r eax
eax                0x3b 59
(gdb) s

Breakpoint 2, _start () at incDec.asm:9
9          inc rax
(gdb) i r eax
eax                0x3c 60
(gdb) s

Breakpoint 3, _start () at incDec.asm:10
10         inc rax
(gdb) i r eax
eax                0x3d 61
(gdb) s

Breakpoint 4, _start () at incDec.asm:11
11         dec rax
(gdb) i r eax
eax                0x3e 62
(gdb) s

Breakpoint 5, _start () at incDec.asm:12
12         dec rax
(gdb) i r eax
eax                0x3d 61
(gdb) s

Breakpoint 6, _start () at incDec.asm:14
14         mov rdi, 0
(gdb) i r eax
eax                0x3c 60
(gdb) s
15         syscall
(gdb) s
[Inferior 1 (process 13422) exited normally]
(gdb)
```



Here **EAX**, which also contains the SysCall ID, is already set to 60, so we can invoke directly the API.

We have shown the **INC** and **DEC** instructions that take as operand a register, and respectively add and subtract 1 from it. The state of **CF** is always preserved.

---

## Integer multiplication

Assembly implements two kind of multiplications, signed and unsigned. Let us begin by the unsigned multiplication, which in its simplest form has the syntax:

**mul <src> <fct>**

in which case, an 'A' register (**AL**, **AX**, **EAX**, **RAX**) must be used for the operands.

The results follow a general rule: if the operands are  $n$  bits, the result will be  $2n$  bits. It follows that the result may need two registers to be stored, in fact in the single operand case the results are stored in the **A** and the **D** registers, as follows:

| First operator | Second operator         | First part of the result | Second part of the result |
|----------------|-------------------------|--------------------------|---------------------------|
| <b>AL</b>      | <b>8 bits operator</b>  | <b>AH</b>                | <b>AL</b>                 |
| <b>AX</b>      | <b>16 bits operator</b> | <b>DX</b>                | <b>AX</b>                 |
| <b>EAX</b>     | <b>32 bits operator</b> | <b>EDX</b>               | <b>EAX</b>                |
| <b>RAX</b>     | <b>64 bits operator</b> | <b>RDX</b>               | <b>RAX</b>                |

Consider the following piece of code:

```
Dump of assembler code for function _start:
0x0000000000400080 <+0>:      mov     al,0x33
0x0000000000400082 <+2>:      mov     bl,0x32
0x0000000000400084 <+4>:      mul     bl
0x0000000000400086 <+6>:      mov     eax,0x3c
0x000000000040008b <+11>:     mov     edi,0x0
0x0000000000400090 <+16>:     syscall
```

We are multiplying two 8-bits values so the result is potentially a 16-bits value.

```

Breakpoint 1, _start () at mul.asm:6
6      mov     al, '3'
(gdb) s
8      mov     bl, '2'
(gdb) s
9      mul     bl
(gdb) i r al
al      0x33 51
(gdb) i r bl
bl      0x32 50
(gdb) s
11     mov     rax, 0x3C

```

Something is already strange. The registers have been loaded with `0x33` and `0x32`, respectively. These are the ASCII codes of 3 and 2, respectively, so something did not work out properly - we will come back on this later, anyway. Since both 51 and 50 are numbers that can be represented with a byte, the multiplication should work as previously described. In fact, EAX contains the correct result of the operation:

```

(gdb) i r bl
bl      0x32 50
(gdb) i r al
al      0xf6 -10
(gdb) i r eax
eax      0x9f6      2550

```

The problem of the 'wrong' values was that we assigned the values of the chars '2' and '3' (note the hyphens).

How can we convert these chars into integers? The next example gives an idea:

```

Dump of assembler code for function _start:
0x0000000000400080 <+0>:      mov     al,0x33 ; '3'
0x0000000000400082 <+2>:      sub     al,0x30 ; '0'...
0x0000000000400084 <+4>:      mov     bl,0x32
0x0000000000400086 <+6>:      sub     bl,0x30
0x0000000000400089 <+9>:      mul     bl
0x000000000040008b <+11>:     mov     eax,0x3c
0x0000000000400090 <+16>:     mov     edi,0x0
0x0000000000400095 <+21>:     syscall

```

## Integer division (`div`)

If  $b \neq 0$ , the division between two integers  $a$  (the dividend) and  $b$  (the divisor) is defined as:

$$a/b = q + r$$

where  $q$  is the quotient and  $r$  is the remainder. These values may also be null.

| Dividend and Divisor size (bytes) | Functioning  |
|-----------------------------------|--|
| <b>2 and 1</b>                    | It is assumed that the dividend is in <b>AX</b> .<br>Quotient is stored in <b>AL</b><br>Remainder is stored in <b>AH</b> |

| Dividend and Divisor size (bytes) | Functioning   |
|-----------------------------------|---|
| <b>4 and 2</b>                    | It is assumed that the dividend is in <b>DX</b> and in <b>AX</b> .<br>Quotient is stored in <b>AX</b><br>Remainder is stored in <b>DX</b>     |
| <b>8 and 4</b>                    | It is assumed that the dividend is in <b>EDX</b> and in <b>EAX</b> .<br>Quotient is stored in <b>EAX</b><br>Remainder is stored in <b>EDX</b> |

The example is not so different from the one regarding multiplication:

```
section .text
global _start

_start:

    mov ax, 23

    mov bl, 8
    div bl

    mov rax, 0x3C
    mov rdi, 0
    syscall

section .data
```

so is the debugging:

```
6          mov ax, 23
(gdb) s
8          mov bl, 8
(gdb) s
9          div bl
(gdb) i r ax
ax          0x17 23
(gdb) i r al
al          0x17 23
(gdb) i r bx
bx          0x8 8
(gdb) p $eflags
$1 = [ IF ]
(gdb) s
11         mov rax, 0x3C
(gdb) s
11         mov rax, 0x3C
(gdb) i r al
al          0x2 2
(gdb) i r ah
ah          0x7 7
(gdb) p $eflags
$2 = [ IF ]
```

The result is coherent: the result (2, in **AL**) times the divisor (8, in **BL**) equals 16, and 16 plus the remainder (7, in **AH**) equals 23 (the dividend, in **AX**).

With this knowledge under our belt, we can start talking about Loops.

# Controlling the program flow

## Loops

---

### The for loop

Loops, in assembly, are implemented by using the instruction '**loop**', unsurprisingly. The instruction provides a simple way to repeat a block of statements a specific number of times. **ECX** is automatically used as a counter and is decremented each time the loop repeats.

The syntax is quite easy:

**loop** <label>

being the argument a text label in the line of code. The execution of the statement involves two steps:

1. Subtracts 1 from **ECX**
2. Compares **ECX** to 0:
  1. If the comparison fails, the program jumps back to <label>
  2. If the comparison succeeds, the program flows passes to the instruction following the **loop** instruction.

An example helps understanding better:

```
global _start
section .text
_start:
    mov rcx, 10
    mov rax, 1

lbl:
    inc rax
    loop lbl

    mov rax, 60
    mov rdi, 0
    syscall

section .data
```

Disassembling may seem tricky

```
(gdb) disassemble _start
Dump of assembler code for function _start:
   0x0000000000400080 <+0>:    mov     ecx,0xa
   0x0000000000400085 <+5>:    mov     eax,0x1
End of assembler dump.
```

This is because we declared a label, thus:

```
(gdb) disassemble lbl
Dump of assembler code for function lbl:
    0x000000000040008a <+0>:      inc    rax
    0x000000000040008d <+3>:      loop   0x40008a <lbl>
    0x000000000040008f <+5>:      mov    eax,0x3c
    0x0000000000400094 <+10>:     mov    edi,0x0
    0x0000000000400099 <+15>:     syscall
End of assembler dump.
```

Setting some breakpoints:

```
(gdb) info breakpoints
Num      Type          Disp Enb Address                What
7        breakpoint     keep y  0x0000000000400080 loop.asm:4
8        breakpoint     keep y  0x0000000000400085 loop.asm:5
9        breakpoint     keep y  0x000000000040008a loop.asm:8
10       breakpoint     keep y  0x000000000040008d loop.asm:9
11       breakpoint     keep y  0x000000000040008f loop.asm:10
```

As usual, we debug the application. The registers are initialised with zeros, as usual:

```
(gdb) i r ecx
ecx      0x0  0
(gdb) i r eax
eax      0x0  0
```

but when entering the loop, these are properly initialised:

```
(gdb) i r ecx
ecx      0xa  10
(gdb) i r eax
eax      0x1  1
```

The body of the loop is constituted by the `inc` and the `loop` statements, so we will just report the steps

| End of step # | i r eax |        | i r ecx |       |
|---------------|---------|--------|---------|-------|
| <b>1</b>      | eax     | 0x2 2  | ecx     | 0x9 9 |
| <b>2</b>      | eax     | 0x3 3  | ecx     | 0x8 8 |
| <b>3</b>      | eax     | 0x4 4  | ecx     | 0x7 7 |
| <b>4</b>      | eax     | 0x5 5  | ecx     | 0x6 6 |
| ...           | ...     |        | ...     |       |
| <b>9</b>      | eax     | 0xa 10 | ecx     | 0x1 1 |
| <b>10</b>     | eax     | 0xa 11 | ecx     | 0x1 0 |

After this last iteration, the program flows returns and the `exit` syscall is invoked.

---

## Infinite loops

Never modify the value of the RCX value during a loop, it could have unpredictable effects. In the next example, we show some code that leads to a never-ending loop ("infinite loop").

```
_start:
    mov rcx, [threshold]

infiniteLoop:
    mov r10, 1
    inc rcx
    loop infiniteLoop

    mov rax, [sys_call]
    mov rdi, [exit_code]
    syscall

section .data
    threshold dq 3
    exit_code dq 0
    sys_call dq 60
```

The reason being the fact that the counter register **RCX** gets incremented in every cycle, thus it will never reach 0.

In fact:

```
Breakpoint 1, infiniteLoop () at buggy.asm:8
8      mov r10, 1
(gdb) i r rcx
rcx      0x3  3
(gdb) s
9      inc rcx
(gdb) i r rcx
rcx      0x3  3
(gdb) s
10     loop infiniteLoop
(gdb) i r rcx
rcx      0x4  4
(gdb) s

Breakpoint 1, infiniteLoop () at buggy.asm:8
8      mov r10, 1
(gdb) i r rcx
rcx      0x3  3
(gdb) ...
```

Beware that also syscalls may change the value of the counter register. In fact, let's get back to our old Hello World:

```

global _start

section .text

_start:
    mov rax, 1
    mov rdi, 1
    mov rsi, hello_world
    mov rdx, length
    syscall
    mov rax, 0x3C
    syscall

section .data
    hello_world: db 'hello world',0xa
    length: equ $-hello_world

```

Let's set some breakpoints and see registers' contents:

```

(gdb) disassemble _start
Dump of assembler code for function _start:
0x0000000004000b0 <+0>:      mov     eax,0x1
0x0000000004000b5 <+5>:      mov     edi,0x1
0x0000000004000ba <+10>:     movabs  rsi,0x6000d4
0x0000000004000c4 <+20>:     mov     edx,0xc
0x0000000004000c9 <+25>:     syscall
0x0000000004000cb <+27>:     mov     eax,0x3c
0x0000000004000d0 <+32>:     syscall
End of assembler dump.
(gdb) break 10
Breakpoint 5 at 0x4000c9: file hello2.asm, line 10.
(gdb) break 11
Breakpoint 6 at 0x4000cb: file hello2.asm, line 11.
...
Breakpoint 5, _start () at hello2.asm:10
10      syscall
(gdb) i r
rax      0x1  1
rbx      0x0  0
rcx      0x0  0
rdx      0xc  12
rsi      0x6000d4  6291668
rdi      0x1  1
rbp      0x0  0x0
rsp      0x7fffffff450  0x7fffffff450
... Breakpoint 6, _start () at hello2.asm:11
11      mov rax, 0x3C
(gdb) i r
rax      0xc  12
rbx      0x0  0
rcx      0x4000b4  4194484
rdx      0xc  12
rsi      0x6000d4  6291668
rdi      0x1  1
rbp      0x0  0x0
rsp      0x7fffffff450  0x7fffffff450
...

```

The debug clearly shows how the API has changed RCX register's contents.

## Putting it all together 1: Fibonacci sequence

The Fibonacci sequence is recursively defined as follows:

$$a_n := \begin{cases} 1 & n = 1 \\ 1 & n = 2 \\ a_{n-1} + a_{n-2} & n > 2 \end{cases}$$

Its elements are 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89, and so on. We wanted to write an assembly program that calculates the 8th term of this sequence (which is 55).

We have written the following code:

```
global _start
section .text
_start:

    mov rax, 1
    mov rbx, 1
    mov rcx, 8
    mov r10, 0

fiboloop:
    mov r10, rbx
    add rbx, rax
    mov rax, r10
    loop fiboloop

    mov rax, 60
    mov rdi, 0
    syscall

section .data
```

The algorithm is quite straightforward: the start section initialises to 1 **RBX** and **RAX** (which, in our case, will always be respectively  $a_{n-1}$  and  $a_{n-2}$ , and a buffer variable (**R10**) that will be used to save the value of **RBX** when swapping variables.

We then start a loop in which:

1. We save the value of **RBX** in **R10**
2. We increase the value of **RBX** with the contents of **RAX**
3. We update the value of **RAX**, putting into it the old **RBX** value (now loaded in **R10**)

It is immediate noticing that this algorithm converges to the intended result.

Debugging this is quite funny. Here we give only the first and the last iterations:



```

(gdb) disassemble _start
Dump of assembler code for function _start:
    0x000000000400080 <+0>:      mov     eax,0x1
    0x000000000400085 <+5>:      mov     ebx,0x1
    0x00000000040008a <+10>:     mov     ecx,0x8
    0x00000000040008f <+15>:     mov     r10d,0x0
End of assembler dump.
(gdb) disassemble fiboloop
Dump of assembler code for function fiboloop:
    0x000000000400095 <+0>:      mov     r10,rbx
    0x000000000400098 <+3>:      add     rbx,rax
    0x00000000040009b <+6>:      mov     rax,r10
    0x00000000040009e <+9>:      loop   0x400095 <fiboloop>
    0x0000000004000a0 <+11>:     mov     eax,0x3c
    0x0000000004000a5 <+16>:     mov     edi,0x0
    0x0000000004000aa <+21>:     syscall
End of assembler dump.

(gdb) i breakpoints
Num      Type      Disp Enb Address      What
5         breakpoint keep y 0x000000000400095 fibonacci.asm:11
6         breakpoint keep y 0x00000000040009e fibonacci.asm:14
7         breakpoint keep y 0x0000000004000a0 fibonacci.asm:15

(gdb) i r
rax      0x1 1
rbx      0x1 1
rcx      0x8 8
rdx      0x0 0
rsi      0x0 0
rdi      0x0 0
rbp      0x0 0x0
rsp      0x7fffffff450 0x7fffffff450
r8       0x0 0
r9       0x0 0
r10      0x0 0
...
rip      0x400095 0x400095 <fiboloop>
eflags   0x202 [ IF ]
cs       0x33 51
ss       0x2b 43
ds       0x0 0
es       0x0 0
fs       0x0 0
gs       0x0 0
...
(gdb) i r
rax      0x15 21
rbx      0x22 34
rcx      0x1 1
rdx      0x0 0
rsi      0x0 0
rdi      0x0 0
rbp      0x0 0x0
rsp      0x7fffffff450 0x7fffffff450
r8       0x0 0
r9       0x0 0
r10      0x15 21
...
rip      0x400095 0x400095 <fiboloop>
eflags   0x216 [ PF AF IF ]
cs       0x33 51
ss       0x2b 43
ds       0x0 0
es       0x0 0
fs       0x0 0
gs       0x0 0

```

This code is not perfect, at least in terms of flexibility (we could have added a threshold constant, for instant, to make it more maintainable), but shows how easy could be solving a problem in assembly.

This can be solved rewriting the code as follows (we also had fun with the syscalls). Up to the reader debugging it.

```
global _start
section .text
_start:

    mov rax, 1
    mov rbx, 1
    mov rcx, [threshold]
    mov r10, 0

fiboloop:
    mov r10, rbx
    add rbx, rax
    mov rax, r10
    loop fiboloop

    mov rax, [sys_call]
    mov rdi, [exit_code]
    syscall

section .data
    threshold dq 8
    exit_code dq 0
    sys_call dq 60
```

## Jumping around the code

A jump is basically a way to skip some parts of the code, or going back in the execution.

---

### Unconditional jump

The simplest form of jump is called unconditional jump, it executes no matter what are the side conditions. The syntax is

`jmp <label>`

the effect is executing the first statement after the textual `label`.

A simple example follows:

```
global _start
section .text
_start:
    mov rax, 0x3C                ;we already prepare the exit syscall
    jmp exitLabel
    mov rdi, 123                 ;this instruction won't ever be executed
exitLabel:
    mov rdi, 0
    syscall

section .data
```

You may have notice the presence of some plain English text after the semicolon. These are comments, and never get considered by the compiler.

The flow of the program is as follows:

```
Dump of assembler code for function _start:
0x000000000400080 <+0>:      mov     eax,0x3c
0x000000000400085 <+5>:      jmp     0x40008c <exitLabel>
0x000000000400087 <+7>:      mov     edi,0x7b
End of assembler dump.
(gdb) disassemble exitLabel
Dump of assembler code for function exitLabel:
0x00000000040008c <+0>:      mov     edi,0x0
0x000000000400091 <+5>:      syscall
End of assembler dump.
(gdb) break _start
Breakpoint 1 at 0x400080: file jump1.asm, line 6.
(gdb) r
Starting program: /home/gbiondo/Development/LX_012_jumps/jump1

Breakpoint 1, _start () at jump1.asm:6
6      mov rax, 0x3C      ;we already prepare the exit syscall
(gdb) s
8      jmp exitLabel
(gdb) i r edi
edi      0x0  0
(gdb) s
exitLabel () at jump1.asm:13
13     mov rdi, 0
(gdb) ...
```

It is evident how the instruction located at <+7> does not get executed. We can cross check that the exit value is, obviously, 0:

```
gbiondo@LinuxMint:~/Development/LX_012_jumps$ ./jump1
gbiondo@LinuxMint:~/Development/LX_012_jumps$ echo $?
0
```

---

## Jump if below (jb)

This is the first form of conditional jumps. It makes the control flow jumping if  $CF = 1$ .

Used after a CMP or SUB instruction, JB transfers control to the label if the first operand was less than the second. (Both operands are treated as unsigned numbers).

Consider the following example:

```

global _start

section .text

_start:
    mov rax, 0x3C      ;we already prepare the exit syscall
    mov rdi, 0         ;initializing the exit code
    jb exitLabel       ;this won't be executed...
    mov rdi, 123        ;changing the exit code
    stc                ;we set the carry flag
    jb exitLabel

exitLabel:
    syscall

section .data

```

Once again, disassembling shows how the program flow works:

```

Breakpoint 1, _start () at jump2.asm:6
6      mov rax, 0x3C      ;we already prepare the exit syscall
(gdb) s
7      mov rdi, 0         ;initializing the exit code
(gdb) s
8      jb exitLabel       ;this won't be executed...
(gdb) s
9      mov rdi, 123        ;changing the exit code
(gdb) p $eflags
$1 = [ IF ]
(gdb) s
10     stc                ;we set the carry flag
(gdb) p $eflags
$2 = [ IF ]
(gdb) s
11     jb exitLabel
(gdb) p $eflags
$3 = [ CF IF ]
(gdb) s
exitLabel () at jump2.asm:14
14     syscall
(gdb) p $eflags
$4 = [ CF IF ]
(gdb) s
[Inferior 1 (process 8253) exited with code 0173]
(gdb) q
gbiondo@LinuxMint:~/Development/LX_012_jumps$ ./jump2
gbiondo@LinuxMint:~/Development/LX_012_jumps$ echo $?
123

```

## Jump if not sign (jns)

This statement performs the jump to the label if and only if the sign flag (SF) is null. The example clarifies.

```

global _start
section .text
_start:
    mov al, 0x2
    sub al, 0x1

land2:
    jns land1
    mov rdi, 0x123

land1:
    sub al, 0x2
    jns land2

    mov rax, 60
    mov rdi, 10
    syscall

section .data

```

Below, the program flow obtained with the debugger:

```

Breakpoint 1, _start () at jump3.asm:7
7      mov al, 0x2
(gdb) s
8      sub al, 0x1
(gdb) s
land2 () at jump3.asm:10
10     jns land1
(gdb) i r al
al      0x1 1
(gdb) p $eflags
$1 = [ IF ]
(gdb) s
land1 () at jump3.asm:14
14     sub al, 0x2
(gdb)
15     jns land2
(gdb) p $eflags
$2 = [ CF PF AF SF IF ]
(gdb) i r al
al      0xff -1
(gdb) s
17     mov rax, 60
(gdb)
18     mov rdi, 10
(gdb)
19     syscall

```

As expected, the first jump is executed ( $2 > 1$ , and hence the subtraction does not set the Sign Flag), whereas the second jump is not.

---

## Other conditional jumps

There are several other conditional jumps to take into consideration. We start with Jump if Equal (`jbe`). Consider the following code:

```

global _start
section .text
_start:
    mov rcx, 0x23
    mov rdx, 0x173
    mov rdi, 100

    cmp rcx, rdx
    je first_label

    cmp rdx, rcx
    je second_label

    cmp rcx, 0x23
    je third_label

first_label:
    mov rdi, 1
    jmp vaivia

second_label:
    mov rdi, 2
    jmp vaivia

third_label:
    mov rdi, 3
    jmp vaivia

vaivia:
    mov rax, 60
    syscall

section .data

```

The **CMP** statement compares the value of a register and another or a register and a constant and sets the status flags in the **EFLAGS** register according to the results. The comparison is performed by subtracting the second operand from the first operand and then setting the status flags in the same manner as the **SUB** instruction.

We will follow first the flow of the code, to determine generic rules for all conditional jumps. The first block of operations is assignments and comparing two different numbers. We do not expect that **je** will be executed.

```

Breakpoint 1, _start () at je.asm:7
7      mov rcx, 0x23
(gdb) s
8      mov rdx, 0x173
(gdb) s
9      mov rdi, 100
(gdb) s
11     cmp rcx, rdx
(gdb) s
12     je first_label
(gdb) i r rcx rdx rdi rip
rcx      0x23 35
rdx      0x173 371
rdi      0x64 100
rip      0x400092 0x400092 <_start+18>
(gdb) p $eflags
$1 = [ CF SF IF ]

```

Observe that **CMP** leverages **SUB**, hence the Carry and the Sign flags are set (the first operand is smaller than the second, hence a negative number is returned). Similarly, also the conditional jump in the following block of code won't be executed:

```
(gdb) s
14          cmp rdx, rcx
(gdb) s
15          je second_label
(gdb) i r rcx rdx rdi rip
rcx          0x23 35
rdx          0x173      371
rdi          0x64 100
rip          0x400097    0x400097 <_start+23>
(gdb) p $eflags
$2 = [ PF IF ]
```

(the parity flag is set because 371-35 gives an even result). There is the last part of code:

```
(gdb) s
17          cmp rcx, 0x23
(gdb) s
18          je third_label
(gdb) s
third_label () at je.asm:29
29          mov rdi, 3
(gdb) i r rcx rdx rdi rip
rcx          0x23 35
rdx          0x173      371
rdi          0x64 100
rip          0x4000ad    0x4000ad <third_label>
(gdb) p $eflags
$3 = [ PF ZF IF ]
(gdb)
```

In this case, the jump is executed (see **RIP**'s value), as the difference between the contents of **RCX** and **0x23** is **0** (differently said, they're the same value).

There are several other conditions that can be used. These are listed in the tables below. The reader is suggested to test these conditional jumps using this methodology.

### Arithmetic operations

| Instruction    | Description                         | Flags tested |
|----------------|-------------------------------------|--------------|
| <b>JE/JZ</b>   | Jump Equal or Jump Zero             | ZF           |
| <b>JNE/JNZ</b> | Jump not Equal or Jump Not Zero     | ZF           |
| <b>JG/JNLE</b> | Jump Greater or Jump Not Less/Equal | OF, SF, ZF   |
| <b>JGE/JNL</b> | Jump Greater/Equal or Jump Not Less | OF, SF       |
| <b>JL/JNGE</b> | Jump Less or Jump Not Greater/Equal | OF, SF       |
| <b>JLE/JNG</b> | Jump Less/Equal or Jump Not Greater | OF, SF, ZF   |

## Logical operations

| Instruction    | Description                            | Flags tested |
|----------------|--|--------------|
| <b>JE/JZ</b>   | Jump Equal or Jump Zero                | ZF           |
| <b>JNE/JNZ</b> | Jump not Equal or Jump Not Zero        | ZF           |
| <b>JA/JNBE</b> | Jump Above or Jump Not Below/<br>Equal | CF, ZF       |
| <b>JAE/JNB</b> | Jump Above/Equal or Jump Not<br>Below  | CF           |
| <b>JB/JNAE</b> | Jump Below or Jump Not Above/<br>Equal | CF           |
| <b>JBE/JNA</b> | Jump Below/Equal or Jump Not<br>Above  | AF, CF       |

## Flag-related operations

| Instruction    | Description                       | Flags tested |
|----------------|-----------------------------------|--------------|
| <b>JXCZ</b>    | Jump if CX is Zero                | none         |
| <b>JC</b>      | Jump If Carry                     | CF           |
| <b>JNC</b>     | Jump If No Carry                  | CF           |
| <b>JO</b>      | Jump If Overflow                  | OF           |
| <b>JNO</b>     | Jump If No Overflow               | OF           |
| <b>JP/JPE</b>  | Jump Parity or Jump Parity Even   | PF           |
| <b>JNP/JPO</b> | Jump No Parity or Jump Parity Odd | PF           |
| <b>JS</b>      | Jump Sign (negative value)        | SF           |
| <b>JNS</b>     | Jump No Sign (positive value)     | SF           |

## The Stack

The stack, in Computer Science, is nothing more than an Abstract Data Type (ADT) that collects coherent elements, and supports (only) two operations:

- push operations: an element is added to the collection
- pop operations: the most recently added element is removed from the collection

From the definition, it emerges that the Stack works in the 'LIFO' mode (Last In, First Out).

In the computer implementation, each thread has a reserved region of memory which is referred as its stack, for when a function executes, it may add some data to this ADT, removing when the



function exits. At a minimum, a thread's stack is used to store the location of function calls in order to allow return statements to return to the correct location.

Another relevant aspect is that the stack is an area of memory with a fixed origin and a variable size. Initially the size of the stack is zero. The stack pointer (hardware register), points to the most recently referenced location on the stack. When the stack null size, the pointer obviously points to the origin.

The stack space is located just under the OS kernel space (more on this later), generally opposite the heap area and grows downwards to lower addresses. (it may grow the opposite direction on some other architectures).

Stack area is devoted to storing all the data needed by a function call in a program. Calling a function is the same as pushing the called function execution onto the top of the stack, and once that function completes, the results are returned popping the function off from the stack.

The dataset pushed for function call is named a stack frame, and it contains the following data:

- the arguments (parameter values) passed to the routine
- the return address back to the routine's caller
- space for the local variables of the routine

**push:** data is placed at the location pointed to by the stack pointer, and the address in the stack pointer is adjusted by the size of the data item

**pop (or pull):** data at the current location pointed to by the stack pointer is removed (and usually put in a special register), the stack pointer is adjusted accordingly

The stack is the memory set aside as scratch space for a thread of execution.

When a function is called, a block is reserved on the top of the stack for local variables and some bookkeeping data. When the function returns, the block becomes unused and can be used the next time a function is called. Mainly, the stack stores local data and return addresses used for parameter passing.

Let us see a practical implementation. Consider the below code:

```
Dump of assembler code for function _start:
0x000000000400080 <+0>:      mov     eax,0x12345678
0x000000000400085 <+5>:      push    rax
0x000000000400086 <+6>:      push    0x9
0x000000000400088 <+8>:      push    0x1234
0x00000000040008d <+13>:     pop     rax
0x00000000040008e <+14>:     pop     rcx
0x00000000040008f <+15>:     pop     rbx
0x000000000400090 <+16>:     mov     eax,0x3c
0x000000000400095 <+21>:     mov     edi,0x0
0x00000000040009a <+26>:     syscall
End of assembler dump.
```

We are using the instruction

`x/4xw $sp`

Whose meaning is as follows:

X - examine

4xv - four (4, but it may be more) words (w) of memory above the stack pointer (here, \$sp) in hexadecimal (x).

At the beginning, the situation is as follows

| Registers     | Stack          |            |            |            |            |
|---------------|----------------|------------|------------|------------|------------|
| rax 0         | 0x7fffffff4c0: | 0x00000001 | 0x00000000 | 0xffffe70b | 0x00007fff |
| rax 305419896 | 0x7fffffff4b8: | 0x12345678 | 0x00000000 | 0x00000001 | 0x00000000 |
| ...           | 0x7fffffff4b0: | 0x00000009 | 0x00000000 | 0x12345678 | 0x00000000 |
| ...           | 0x7fffffff4a8: | 0x00001234 | 0x00000000 | 0x00000009 | 0x00000000 |
|               | 0x7fffffff4b8: | 0x12345678 | 0x00000000 | 0x00000001 | 0x00000000 |
| rax 4660      | 0x7fffffff4b0: | 0x00000009 | 0x00000000 | 0x12345678 | 0x00000000 |
|               | 0x7fffffff4c0: | 0x00000001 | 0x00000000 | 0xffffe70b | 0x00007fff |
| rax 4660      | 0x7fffffff4b8: | 0x12345678 | 0x00000000 | 0x00000001 | 0x00000000 |
| rbx 305419896 | 0x7fffffff4c8: | 0xffffe70b | 0x00007fff | 0x00000000 | 0x00000000 |
| rax 4660      | 0x7fffffff4c0: | 0x00000001 | 0x00000000 | 0xffffe70b | 0x00007fff |
| rbx 305419896 | 0x7fffffff4d0: | 0x00000000 | 0x00000000 | 0xffffe72b | 0x00007fff |
| rcx 9         |                |            |            |            |            |
| ...           | ...            |            |            |            |            |

Note how the stack shrunk and enlarged during the program run.

## Logical operations

We are now examining the bitwise logical operations. Assembly on Linux implements the following logical operators:

- AND
- OR
- NOT
- TEST
- XOR

We are going to see how they work shortly. It is beneficial to remember the definitions of these operands:

| A | B | NOT A | A AND B | A OR B | A XOR B |
|---|---|-------|---------|--------|---------|
| 0 | 0 | 1     | 0       | 0      | 0       |
| 0 | 1 | 1     | 0       | 1      | 1       |
| 1 | 0 | 0     | 0       | 1      | 1       |
| 1 | 1 | 0     | 1       | 1      | 0       |

Bitwise operation means that each operand is analysed at a bit level, for instance:

|                  |          |          |          |          |          |          |          |          |
|------------------|----------|----------|----------|----------|----------|----------|----------|----------|
| <b>A</b>         | 0        | 1        | 0        | 1        | 0        | 0        | 1        | 0        |
| <b>B</b>         | 1        | 1        | 1        | 0        | 0        | 1        | 1        | 1        |
| <b>A   B</b>     | <b>1</b> | <b>1</b> | <b>1</b> | <b>1</b> | <b>0</b> | <b>1</b> | <b>1</b> | <b>1</b> |
| <b>A &amp; B</b> | <b>0</b> | <b>1</b> | <b>0</b> | <b>0</b> | <b>0</b> | <b>0</b> | <b>1</b> | <b>0</b> |

Let's dive down to the code.

In this exercise we will use the values in the table below; we are also computing the results manually to cross check.

|                    |          |          |          |          |          |          |          |          |
|--------------------|----------|----------|----------|----------|----------|----------|----------|----------|
| <b>RAX</b>         | 1        | 0        | 1        | 1        | 1        | 0        | 1        | 1        |
| <b>RBX</b>         | 1        | 1        | 0        | 1        | 0        | 1        | 1        | 0        |
| <b>RAX AND RBX</b> | <b>1</b> | <b>0</b> | <b>0</b> | <b>1</b> | <b>0</b> | <b>0</b> | <b>1</b> | <b>0</b> |
| <b>RAX OR RBX</b>  | <b>1</b> | <b>1</b> | <b>1</b> | <b>1</b> | <b>1</b> | <b>0</b> | <b>1</b> | <b>1</b> |
| <b>RAX XOR RBX</b> | <b>0</b> | <b>1</b> | <b>1</b> | <b>0</b> | <b>1</b> | <b>1</b> | <b>0</b> | <b>1</b> |
| <b>NOT RAX</b>     | <b>0</b> | <b>1</b> | <b>0</b> | <b>0</b> | <b>0</b> | <b>1</b> | <b>0</b> | <b>0</b> |

Consider then this code:

```
Dump of assembler code for function _start:
0x000000000400080 <+0>:      mov     eax,0x10111011
0x000000000400085 <+5>:      mov     ebx,0x11010110
0x00000000040008a <+10>:     and     rax,rbx
0x00000000040008d <+13>:     mov     eax,0x10111011
0x000000000400092 <+18>:     mov     ebx,0x11010110
0x000000000400097 <+23>:     or      rax,rbx
0x00000000040009a <+26>:     mov     eax,0x10111011
0x00000000040009f <+31>:     mov     ebx,0x11010110
0x0000000004000a4 <+36>:     xor     rax,rbx
0x0000000004000a7 <+39>:     mov     eax,0x10111011
0x0000000004000ac <+44>:     not     rax
0x0000000004000af <+47>:     mov     eax,0x10111011
0x0000000004000b4 <+52>:     mov     ebx,0x11010110
0x0000000004000b9 <+57>:     test    rax,rbx
0x0000000004000bc <+60>:     mov     eax,0x3c
0x0000000004000c1 <+65>:     mov     edi,0xa
0x0000000004000c6 <+70>:     syscall
```

The result of the AND operation is stored again in RAX, as shown below:

```

10          and rax, rbx
(gdb) i r rax rbx rip eflags
rax          0x10111011  269553681
rbx          0x11010110  285278480
rip          0x40008a    0x40008a <_start+10>
eflags      0x202      [ IF ]
(gdb) s
12          mov rax,0x10111011
(gdb) i r rax rbx rip eflags
rax          0x10010010  268501008
rbx          0x11010110  285278480
rip          0x40008d    0x40008d <_start+13>
eflags      0x202      [ IF ]

```

The second operation behaves in the same way.

```

15          or rax, rbx
(gdb) i r rax rbx rip eflags
rax          0x10111011  269553681
rbx          0x11010110  285278480
rip          0x400097    0x400097 <_start+23>
eflags      0x202      [ IF ]
(gdb) s
17          mov rax,0x10111011
(gdb) i r rax rbx rip eflags
rax          0x11111111  286331153
rbx          0x11010110  285278480
rip          0x40009a    0x40009a <_start+26>
eflags      0x206      [ PF IF ]

```

As for the **XOR**, once again the result is stored in **RAX**:

```

(gdb) i r rax rbx rip eflags
rax          0x10111011  269553681
rbx          0x11010110  285278480
rip          0x4000a4    0x4000a4 <_start+36>
eflags      0x206      [ PF IF ]
(gdb) s
22          mov rax,0x10111011
(gdb) i r rax rbx rip eflags
rax          0x1101101  17830145
rbx          0x11010110  285278480
rip          0x4000a7    0x4000a7 <_start+39>
eflags      0x202      [ IF ]

```

The **NOT** operator behaves in the same manner:

```

23          not rax
(gdb) i r rax rbx rip eflags
rax          0x10111011  269553681
rbx          0x11010110  285278480
rip          0x4000ac    0x4000ac <_start+44>
eflags      0x202      [ IF ]
(gdb) s
25          mov rax,0x10111011
(gdb) i r rax rbx rip eflags
rax          0xfffffffffefeefee-269553682
rbx          0x11010110  285278480
rip          0x4000af    0x4000af <_start+47>
eflags      0x202      [ IF ]

```

The **TEST** operator, finally, behaves like **AND** and it is not discussed here.

## Subroutines

Subroutine definition follows this syntax:

```
subroutine_name:
    <statements>
    ret
```

Subroutines must be defined just after the `_start` section, they begin with a symbolic name, and end with the `ret` statement.

Consider the following piece of code:

```
global _start

section .text

addition:

    mov eax, ecx
    add eax, edx
    ret

_start:

    mov ecx, 0x14
    mov edx, 0x15
    call addition

    mov r8, 0x4
    mov r9, 0x2
    call addition

    mov rax, 60
    mov rdi, 1
    syscall

section .data
```

Into GDB we have:

```
Dump of assembler code for function _start:
0x000000000400085 <+0>:      mov     ecx,0x14
0x00000000040008a <+5>:      mov     edx,0x15
0x00000000040008f <+10>:     call    0x400080 <addition>
0x000000000400094 <+15>:     mov     r8d,0x4
0x00000000040009a <+21>:     mov     r9d,0x2
0x0000000004000a0 <+27>:     call    0x400080 <addition>
0x0000000004000a5 <+32>:     mov     eax,0x3c
0x0000000004000aa <+37>:     mov     edi,0x1
0x0000000004000af <+42>:     syscall
End of assembler dump.
(gdb) disassemble addition
Dump of assembler code for function addition:
0x000000000400080 <+0>:      mov     eax,ecx
0x000000000400082 <+2>:      add     eax,edx
0x000000000400084 <+4>:      ret
End of assembler dump.
```

It is important that the reader now puts attention also on statements' addresses, since we will take into consideration also the **RSP** register, which points to the address of the next instruction to be executed.

```
Breakpoint 1, _start () at subs.asm:13
13      mov ecx, 0x14
(gdb) i r
...
rcx          0x14 20
...
rsp          0x7fffffff480 0x7fffffff480
...rip       0x400085 0x400085 <_start>
eflags       0x202 [ IF ]
...
```

This is correct, since the first instruction has not been run yet. Observe that first instruction to be ran is the first after the **RET**.

```

14      mov edx, 0x15
(gdb) i r
...
rcx      0x14 20
...
rsp      0x7fffffff480    0x7fffffff480
...
rip      0x40008a    0x40008a <_start+5>
eflags   0x202    [ IF ]
...
(gdb) s
15      call addition
(gdb) info registers
...
rcx      0x14 20
rdx      0x15 21
...
rsp      0x7fffffff480    0x7fffffff480
...
rip      0x40008f    0x40008f <_start+10>
eflags   0x202    [ IF ]
...
(gdb) s
addition () at subs.asm:7
7      mov eax, ecx
(gdb) info registers
...
rcx      0x14 20
rdx      0x15 21
...
rsp      0x7fffffff478    0x7fffffff478
...
rip      0x400080    0x400080 <addition>
eflags   0x202    [ IF ]
...
(gdb) s
8      add eax, edx
(gdb) i r
rax      0x14 20
rbx      0x0 0
rcx      0x14 20
rdx      0x15 21
...
rsp      0x7fffffff478    0x7fffffff478
...
rip      0x400082    0x400082 <addition+2>
eflags   0x202    [ IF ]
...

```

```

(gdb) s
addition () at subs.asm:9
9          ret
(gdb) i registers
rax          0x29 41
rbx          0x0  0
rcx          0x14 20
rdx          0x15 21
...
rsp          0x7fffffff478    0x7fffffff478
...
rip          0x400084    0x400084 <addition+4>
eflags      0x202      [ IF ]
...
17          mov r8,0x4
(gdb) i r
rax          0x29 41
rbx          0x0  0
rcx          0x14 20
rdx          0x15 21
...
rsp          0x7fffffff480    0x7fffffff480
...
rip          0x400094    0x400094 <_start+15>
eflags      0x202      [ IF ]
...

```

Observe the following:

1. The next instruction pointed is **<\_start+15>**, which is the first instruction after the subroutine invocation
2. The result is stored in **RAX** (more properly, in **EAX**).

Now the remaining part of the program is straightforward, but what happens in the subroutine the second time? We just ran it, and we skip debugging, obtaining:

```

(gdb) i r
rax          0x29 41
rbx          0x0  0
rcx          0x14 20
rdx          0x15 21
rsi          0x0  0
rdi          0x0  0
rbp          0x0  0x0
rsp          0x7fffffff480    0x7fffffff480
r8           0x4  4
r9           0x2  2
r10          0x0  0
r11          0x0  0
r12          0x0  0
r13          0x0  0
r14          0x0  0
r15          0x0  0
rip          0x4000a5    0x4000a5 <_start+32>
eflags      0x202      [ IF ]
cs          0x33 51
ss          0x2b 43
ds          0x0  0
es          0x0  0
fs          0x0  0
gs          0x0  0

```

The computation still took place on the first given registers, which does not add flexibility to the subroutine mechanism. We will now see another, more flexible version.



